



## Exercise 2.3: Color

1. Load the image file `girl.tif`. Check the size of array for this image by using the command `print(image.shape)`, where `image` is the image matrix. Also, print the data type of this matrix.


Notice that this is a three dimensional array of type `uint8`. It contains three gray scale image planes corresponding to the red, green, and blue components for each pixel. Since each color pixel is represented by three bytes, this is commonly known as a 24-bit image.

```
In [155]: girl = plt.imread("girl.tif")  
print(girl.shape)
```

```
(256, 384, 3)
```

2. Display the image. Note that `cmap`, `vmin`, `vmax` arguments are not needed.

```
In [156]: plt.imshow(girl)  
plt.show()
```



3. Extract each of the color components, then plot each color component.

Note that while the original is a color image, each color component separately is a monochrome image, so plotting each color component requires `cmap`, `vmin`, `vmax`

```
In [157]: plt.imshow(girl[:, :, 0], cmap='gray', vmin=0, vmax=255)  
plt.show()
```



```
In [158]: plt.imshow(girl[:, :, 1], cmap='gray', vmin=0, vmax=255)  
plt.show()
```



```
In [159]: plt.imshow(girl[:, :, 2], cmap='gray', vmin=0, vmax=255)  
plt.show()
```



**4. Load the files `ycbcr.npy` using `np.load()`**

**(<https://numpy.org/doc/stable/reference/generated/numpy.load.html>), and print its type and data shape dtype .**

This file contains a NumPy array for a color image in  $YC_bC_r$  format. The array contains three gray scale image planes that correspond to the luminance ( $Y$ ) and two chrominance ( $C_bC_r$ ) components.

```
In [160]: ycbcr = np.load("ycbcr.npy")
print(ycbcr.dtype)
print(ycbcr.shape)
B = ycbcr.astype(float)
```


```
uint8
(256, 384, 3)
```

**5. Plot each of the components.**


```
In [161]: plt.imshow(ycbcr[:, :, 0], cmap='gray', vmin=0, vmax=255)
plt.show()
```



```
In [162]: plt.imshow(ycbcr[:, :, 1], cmap='gray', vmin=0, vmax=255)  
plt.show()
```



```
In [163]: plt.imshow(ycbcr[:, :, 2], cmap='gray', vmin=0, vmax=255)  
plt.show()
```



**6. Complete the function below that will perform the transformation of equation (2). It should accept a 3-D  $YC_bC_r$  image array as input, and return a 3-D  $RGB$  image array.**

- Make sure `ycbcr` is in `double` or `float` before any processing.
- After conversion, to make sure the values of `rgb` are in `[0, 255]`, use `np.clip()`. (<https://numpy.org/doc/stable/reference/generated/numpy.clip.html>).

```
In [164]: def ybcr2rgb(ycbcr):
    """
    Parameters
    ---
    ybcr: image in YCbCr

    Returns
    ---
    rgb: image RGB
    """

    a = np.zeros(ycbcr.shape)
    rgb = np.zeros(ycbcr.shape)
    a[:, :, 0] = ybcr[:, :, 0] + 1.4025*(ycbcr[:, :, 2] - 128)
    a[:, :, 1] = ybcr[:, :, 0] - 0.3443*(ycbcr[:, :, 1] - 128) - 0.7144*(ycbcr[:, :, 2] - 128)
    a[:, :, 2] = ybcr[:, :, 0] + 1.7730*(ycbcr[:, :, 1] - 128)
    np.clip(a, 0, 255, out=rgb)
    return rgb
```

**7. Now, convert the `ycbcr` array to an `RGB` representation and display the color image.**

- Before displaying the image, make sure its data type is `np.uint8`.



```
In [165]: rgb = ycbcr2rgb(B)
print(rgb.dtype)

plt.imshow(rgb.astype(np.uint8))
plt.show()
```

float64



**8. Load the file `h.npy` . This is a  $5 \times 5$  Gaussian filter with  $\sigma^2 = 2.0$ . (See the first week of the experiment for more details on this type of filter.)**

```
In [166]: h = np.load("h.npy")
```

**9. Alter the `ycbcr` array by filtering only the luminance component, `ycbcr[:, :, 0]` , using the Gaussian filter (use `convolve2d()` function from last lab). Convert the result to RGB, and display it.**

- Instead of altering the original `ycbcr` , you can create a copy by `ycbcr1 = ycbcr.copy()` .



```
In [167]: def convolve2d(image, kernel):
          """
          Parameters
          ---
          image: the input image
          kernel: the filter

          Returns
          ---
          filtered: the filtered image
          """
          H,W = np.shape(image)
          s = len(kernel)
          filtered = np.zeros((H,W))
          new = np.zeros((H+s-1,W+s-1))

          for i in range(H):
              for j in range(W):
                  new[i+s//2,j+s//2] = image[i,j]

          h = np.flip(kernel)


          for k in range(H):
              for l in range(W):
                  filtered[k,l] = np.sum((new[k:k+s,l:l+s]*h))
          return filtered
```

```
In [168]: ybcr1 = ybcr.copy()

ybcr1[:, :, 0] = convolve2d(ybcr1[:, :, 0], h)

rgb1 = ybcr2rgb(ybcr1)

plt.imshow(rgb1.astype(np.uint8))
plt.show()
```




**10. Now alter `ybcr` by filtering both chrominance components, `ybcr[:, :, 1]` and `ybcr[:, :, 2]`, using the Gaussian filter. Convert this result to RGB, and display it.**

- Again, instead of altering the original `ybcr`, you can create a copy by `ybcr2 = ybcr.copy()`.

```
In [169]: ybcr2 = ybcr.copy()

ybcr2[:, :, 1] = convolve2d(ybcr2[:, :, 1], h)
ybcr2[:, :, 2] = convolve2d(ybcr2[:, :, 2], h)
rgb2 = ybcr2rgb(ybcr2)

plt.imshow(rgb2.astype(np.uint8))
plt.show()
```



## Exercise 3.2: Halftoning - Simple Thresholding

1. Load the grayscale image file `house.tif` and display it.

```
In [170]: house = plt.imread("house.tif")  
  
plt.imshow(house, cmap='gray', vmin=0, vmax=255)  
plt.show()
```



**2. Try the simple thresholding technique based on equation (3), using  $T = 108$ , and display the result.**

- In Python, an easy way to threshold an image  $X$  is to use the command  $Y = 255 * (X > T)$ .

```
In [171]: Y = 255*(house>108)

plt.imshow(Y, cmap='gray', vmin=0, vmax=255)
plt.show()
```



**3. Create an “absolute error” image by subtracting the binary from the original image, and then taking the absolute value. The degree to which the original image is present in the error image is a measure of signal dependence of the quantization error. Display the error image.**

```
In [172]: Z = np.abs(house - Y)

plt.imshow(Z, cmap='gray', vmin=0, vmax=255)
plt.show()
```



**4. Compute the mean square error (MSE), which is defined by**

$$\text{MSE} = \frac{1}{MN} \sum_{i,j} \{f[i,j] - b[i,j]\}^2 \quad (9)$$

**where  $MN$  is the total number of pixels in each image,  $f$  is the original image and  $b$  is the binarized image.**

```
In [173]: H,W = np.shape(house)
M = H*W
H1,W1 = np.shape(Y)
N = H1*W1

MSE = np.sum((house-Y)**2)/(M*N)
print(MSE)
```

0.084497663916813

## Exercise 3.4: Halftoning - Ordered Dithering

1. Based on this index matrix and equation (6), create the corresponding threshold matrix and print it.

```
In [174]: I = np.array([[12, 8, 10, 6], [4, 16, 2, 14], [9, 5, 11, 7], [1, 13, 3, 15]])
H,W = np.shape(I)
T = 255 * (I - 0.5)/(H*W)
print(T)
```

```
[[183.28125 119.53125 151.40625  87.65625]
 [ 55.78125 247.03125  23.90625 215.15625]
 [135.46875  71.71875 167.34375 103.59375]
 [  7.96875 199.21875  39.84375 231.09375]]
```

2. Apply the ordered dithering and display the halftoned image.




```
In [175]: H,W = np.shape(house)
num_y = H//4
num_x = W//4

m = np.tile(T, (num_y, num_x))

Y = 255*(house>m)

plt.imshow(Y, cmap='gray', vmin=0, vmax=255)
plt.show()
```



**3. Compute the error image and display it.**

```
In [176]: Z = np.abs(house - Y)

plt.imshow(Z, cmap='gray', vmin=0, vmax=255)
plt.show()
```



#### 4. Compute the MSE of the error image.

```
In [177]: H,W = np.shape(house)
M = H*W
H1,W1 = np.shape(Y)
N = H1*W1

MSE = np.sum((house-Y)**2)/(M*N)
print(MSE)
```

0.1327400335835086

### Exercise 3.6: Halftoning - Error Diffusion

1. Use the algorithm to create the halftoned image and display it.

```
In [181]: house0 = plt.imread("house.tif")
house1 = house0.astype(np.double).copy()
H,W = np.shape(house1)
output = np.zeros((H,W))
e = np.zeros(house1.shape)

for i in range(1,H-1):
    for j in range(1,W-1):
        output[i,j] = 255*(house1[i,j]>108)
        e[i,j] = house1[i,j] - output[i,j]
        output[i,j+1] += e[i,j]*7/16
        output[i+1,j-1] += e[i,j]*3/16
        output[i+1,j] += e[i,j]*5/16
        output[i+1,j+1] += e[i,j]*1/16

plt.imshow(output.astype(np.uint8), cmap='gray', vmin=0, vmax=255)
plt.show()
plt.imshow(e.astype(np.uint8), cmap='gray', vmin=0, vmax=255)
plt.show()
```





**2. Compute the error image and display it.**

In [ ]:

**3. Compute the MSE of the error image.**

In [179]: *# insert your code here*

**4. By comparing three MSE values, is the MSE consistent with the visual quality?**

insert your answer here

**5. By looking at the error images, determine which method appears to be the least signal dependent? Does the signal dependence seem to be correlated with the visual quality?**

insert your answer here

## Exercise 3.7: Halftoning - Filtered Halftone

1. The human visual system naturally lowpass filters halftone images. To analyze this phenomenon, filter each of the halftone images with the Gaussian lowpass filter `h` that you loaded from `h.npy`, and measure the MSE of the filtered versions.

In [180]: *# insert your code here*

2. Use the following template to make a table.

Halftone Method	Filtered	Not filtered
Simple Thresholding		
Ordered Dithering		
Error Diffusion		

3. Compare the MSE's of the filtered versions with the nonfiltered versions for each method. What is the implication of these observations with respect to how we perceive halftone images.

insert your answer here